


[EvoNexus · Oct 8th]

Scale Software Development with AI Sandboxes

From laptop steering to agent swarms

Mathew Mathew · CTO / Claramap.com]





LET'S CONNECT

Mathew Mathew

CTO & Founder, Claramap

Claramap helps companies design, build and run **agentic systems** in production.

20+ years as a software engineer and data scientist at IBM Cloud, Amdocs, AT&T, SolarWinds and Discover.



SCAN TO CONNECT

[linkedin.com/in/mathewma](https://www.linkedin.com/in/mathewma)

github.com/mathaix/fife

github.com/mathaix/signals

BEFORE WE START

Where are you on the ladder?

The diagram illustrates five levels of AI agent integration, arranged vertically. A central navigation bar is positioned between the second and third levels. The levels are numbered 1 to 5 from bottom to top. Each level includes a number, a title, a description, and a key feature or bottleneck.

Level	Agent Type	Description	Key Feature / Bottleneck
5	Agent swarm	Many agents, one problem	Isolation vs. speed vs. cost
4	Agent factory	Tickets in, PRs out	Review & verification
▲ SANDBOXES · · · · · LAPTOP ▼ BOTTLENECK			
3	Parallel local agents	Worktrees and YOLO mode	Your hardware
2	Local agent	An agent in your terminal	Your approvals
1	Autocomplete	AI suggests, you do the rest	Your typing

BOTTLENECK

What a sandbox is

An ephemeral, isolated computer with an API, where an agent can run anything inside a blast radius **you control.**

Isolated

Ephemeral

Programmable

Fast to start

Snapshottable

The isolation spectrum

OS sandbox

Seatbelt,
bubblewrap,
Landlock

Container

Namespaces and
cgroups; shared
host kernel

gVisor

User-space kernel
intercepts syscalls

microVM

Firecracker, Kata;
own kernel

Full VM

Full hardware
virtualization

← FASTER START • LESS OVERHEAD


STRONGER BOUNDARY →

Who builds the sandboxes

	Tool	Isolation	\$/hr*	Notable
	Strands Box PREVIEW	OS sandbox + policy	0SS · your Mac	From AWS, Oct 7; policies see the agent's history
	Docker	Container	0SS · your infra	Dev default; not enough alone for untrusted code
	Daytona	Container (Sysbox)	\$0.08	Fast creation; egress allowlists; placeholder secrets
	Modal	gVisor	\$0.12	Built for scale; locked down by default
	E2B	Firecracker microVM	\$0.08	Agent-focused code-interpreter SDK
	Lambda MicroVMs	Firecracker microVM	\$0.13	Suspend and resume; sessions up to 8h
	Firecracker	microVM	0SS · KVM hosts	Open-source building block from AWS


* Managed: 1 vCPU + 2 GiB for one hour, US list prices, Oct 2026. Open source: no license fee, but you run and pay for the machines. Ordered as on the isolation spectrum: faster start → stronger boundary.

The agent loop



Every action is a tool you defined. You know exactly what it can do.

The coding agent




Give it bash and it writes its own tools. You can't predict what it will run, so run it in a sandbox.

Headless execution

```
$ claude -p "Fix the failing test in auth/" --output-format json
```

Decompose

Split a big task. Each subagent gets one job and only the context it needs.




Automate

No one at the keyboard. Events and schedules start the agent.



Orchestrate



Same agent, many entry points. Every task gets its own sandbox.

Laptops are where agents learn to code.
Sandboxes are where they learn to scale.

Thank you · Questions?

Mathew Mathew · mmathew@claramap.com